

# Regression

Jonathan Ho

September 23, 2022

## Linear Regression

Linear regression is the concept of comparing two variables to observe their relationship. More specifically, the predictors  $x$  will be used to try and predict the target  $y$ . In linear regression, the target would be a quantitative value. Through the use of residuals along side statistical values, an analysis of the linear regression can be carried out to check if it is a good fit given the predictors and target values. Linear regression is relatively simple to understand, is powerful when the data supports a linear relationship, and it has low variance. However, a drawback would be the high bias it has due to the linear shape being assumed with the data.

## Loading in the data

Loading in the csv file.

```
df <- read.csv("vgsales.csv")
str(df)
```

```
## 'data.frame': 16598 obs. of 11 variables:
## $ Rank      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ Name      : chr  "Wii Sports" "Super Mario Bros." "Mario Kart Wii" "Wii Sports Resort"
##           ...
## $ Platform  : chr  "Wii" "NES" "Wii" "Wii" ...
## $ Year      : chr  "2006" "1985" "2008" "2009" ...
## $ Genre     : chr  "Sports" "Platform" "Racing" "Sports" ...
## $ Publisher : chr  "Nintendo" "Nintendo" "Nintendo" "Nintendo" ...
## $ NA_Sales  : num  41.5 29.1 15.8 15.8 11.3 ...
## $ EU_Sales  : num  29.02 3.58 12.88 11.01 8.89 ...
## $ JP_Sales  : num  3.77 6.81 3.79 3.28 10.22 ...
## $ Other_Sales : num  8.46 0.77 3.31 2.96 1 0.58 2.9 2.85 2.26 0.47 ...
## $ Global_Sales: num  82.7 40.2 35.8 33 31.4 ...
```

## Dividing up the data

Splitting the data into 80/20 train/test. There were two rows within the dataset that were corrupted or wrong, so the removal of them was necessary.

```
# Remove two genres that are outliers
new_df <- df[!grepl("Sony Computer Entertainment", df$Genre),]
new_df <- new_df[!grepl("Idea Factory", new_df$Genre),]

set.seed(1234)
x <- sample(1:nrow(new_df), nrow(new_df)*0.8, replace=FALSE)
train <- new_df[x,]
test <- new_df[-x,]
```

## Data exploration

1. See all the columns and what data type they are.

```
str(train)
```

```
## 'data.frame': 13276 obs. of 11 variables:
## $ Rank : int 7453 8017 7163 8087 9197 623 15245 10886 935 12690 ...
## $ Name : chr "Mortal Kombat: Special Forces" "Reel Fishing II" "Dark Souls II" "Battle Commander: Hachibushu Shura no Heihou" ...
## $ Platform : chr "PS" "PS" "XOne" "SNES" ...
## $ Year : chr "2000" "2000" "2015" "1991" ...
## $ Genre : chr "Fighting" "Sports" "Role-Playing" "Strategy" ...
## $ Publisher : chr "Midway Games" "Victor Interactive" "Namco Bandai Games" "Banpresto" ...
## $ NA_Sales : num 0.12 0.1 0.13 0 0.11 1.15 0.02 0.09 0.85 0 ...
## $ EU_Sales : num 0.08 0.07 0.07 0 0.03 1.14 0 0 0.71 0.05 ...
## $ JP_Sales : num 0 0 0.18 0 0.06 0 0 0.13 0 ...
## $ Other_Sales : num 0.01 0.01 0.02 0 0 0.13 0 0.01 0.16 0.01 ...
## $ Global_Sales: num 0.21 0.18 0.22 0.18 0.14 2.48 0.02 0.09 1.86 0.06 ...
```

2. First and last six rows of each column.

```
head(train)
```

R...	Name	Platform	Y...	Genre	
<int>	<chr>	<chr>	<chr>	<chr>	
7452 7453	Mortal Kombat: Special Forces	PS	2000	Fighting	
8016 8017	Reel Fishing II	PS	2000	Sports	
7162 7163	Dark Souls II	XOne	2015	Role-Playing	
8086 8087	Battle Commander: Hachibushu Shura no Heihou	SNES	1991	Strategy	
9196 9197	NFL Blitz 20-03	GC	2002	Sports	
623 623	Tomb Raider: The Last Revelation	PS	1998	Action	
6 rows   1-6 of 12 columns					

```
tail(train)
```

	<b>Rank</b> <int>	<b>Name</b> <chr>	<b>Platform</b> <chr>	<b>Y...</b> <chr>	<b>Genre</b> <chr>	<b>Publisher</b> <chr>
16418	16420	th!nk Logic Trainer	Wii	2009	Puzzle	Conspiracy Entertainment
12963	12964	Cities in Motion	PC	2011	Simulation	Paradox Interactive
13301	13302	Gummy Bears Mini Golf	DS	2010	Sports	Storm City Games
11025	11026	Hello Kitty's Cube Frenzy	PS	1998	Puzzle	Culture Publishers
14726	14728	Auto Modellista	GC	2003	Racing	Capcom
11750	11751	Super Baseball	2600	1987	Sports	Atari

6 rows | 1-8 of 12 columns

3. Checking how many NAs there are.

```
apply(train, function(y) sum(is.na(y)))
```

```
##           Rank           Name           Platform           Year           Genre           Publisher
##           0             0             0             0             0             0
##   NA_Sales   EU_Sales   JP_Sales   Other_Sales   Global_Sales
##           0             0             0             0             0
```

4. Display how many rows and columns there are in the data.

```
dim(train)
```

```
## [1] 13276    11
```

5. Summary of the statistics of each column.

```
summary(train)
```

```
##      Rank      Name      Platform      Year
## Min.    :    1 Length:13276 Length:13276 Length:13276
## 1st Qu.: 4164 Class :character Class :character Class :character
## Median : 8352 Mode  :character Mode  :character Mode  :character
## Mean    : 8326
## 3rd Qu.:12475
## Max.    :16600
##      Genre      Publisher      NA_Sales      EU_Sales
## Length:13276 Length:13276 Min.    : 0.000 Min.    : 0.0000
## Class :character Class :character 1st Qu.: 0.000 1st Qu.: 0.0000
## Mode  :character Mode  :character Median : 0.080 Median : 0.0200
##                                     Mean    : 0.262 Mean    : 0.1448
##                                     3rd Qu.: 0.240 3rd Qu.: 0.1100
##                                     Max.    :41.490 Max.    :29.0200
##      JP_Sales      Other_Sales      Global_Sales
## Min.    : 0.00000 Min.    : 0.00000 Min.    : 0.0100
## 1st Qu.: 0.00000 1st Qu.: 0.00000 1st Qu.: 0.0600
## Median : 0.00000 Median : 0.01000 Median : 0.1700
## Mean    : 0.07772 Mean    : 0.04825 Mean    : 0.5346
## 3rd Qu.: 0.04000 3rd Qu.: 0.03000 3rd Qu.: 0.4700
## Max.    :10.22000 Max.    :10.57000 Max.    :82.7400
```

6. See all the qualitative descriptions in the Genre column for the train and test.

```
table(train$Genre)
```

```
##
##      Action      Adventure      Fighting      Misc      Platform      Puzzle
##      2637      1039      675      1397      705      465
##      Racing Role-Playing      Shooter      Simulation      Sports      Strategy
##      998      1191      1056      708      1858      547
```


```
table(test$Genre)
```

```
##
##      Action      Adventure      Fighting      Misc      Platform      Puzzle
##      679      245      173      342      181      117
##      Racing Role-Playing      Shooter      Simulation      Sports      Strategy
##      251      297      254      159      488      134
```

## Data visualization


Box plot of the Global\_Sales. Within the dataset, there are a good amount of numbers outside the middle half of the sample making it seem like there are a lot of outliers.

```
boxplot(train$Global_Sales, xlab="Global Sales", main="test", ylim=c(0,1.2))
```



A bar plot of Platform to see the frequency of each within the dataset.

```
counts <- table(train$Platform)
barplot(counts, xlab="Platform", ylab="Frequency", col="dodgerblue4", las=2) # las=2 displays all the Platforms
```




A pie chart of Genre to visually see the distribution of game genres.

```
slices <- c(sum(train$Genre=="Action"), sum(train$Genre=="Adventure"), sum(train$Genre=="Fighting"),
sum(train$Genre=="Misc"), sum(train$Genre=="Platform"), sum(train$Genre=="Puzzle"), sum(train$Genre=="Racing"),
sum(train$Genre=="Role-Playing"), sum(train$Genre=="Shooter"), sum(train$Genre=="Simulation"), sum(train$Genre=="Sports"), sum(train$Genre=="Strategy"))

lbls <- c("Action", "Adventure", "Fighting", "Misc", "Platform", "Puzzle", "Racing", "Role-Playing", "Shooter", "Simulation", "Sports", "Strategy")

pie(slices, labels=lbls, main="Game Genres")
```

## Game Genres



## Simple linear regression

A single predictor is used to see the impact of Genre on Global\_Sales.

```
lm1 <- lm(Global_Sales~Genre, data=train)
summary(lm1)
```

```
##
## Call:
## lm(formula = Global_Sales ~ Genre, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.868 -0.462 -0.308 -0.042  82.145
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.53198    0.03031  17.549 < 2e-16 ***
## GenreAdventure -0.33847    0.05702   -5.936 2.99e-09 ***
## GenreFighting  -0.02819    0.06715   -0.420 0.674652
## GenreMisc      -0.07257    0.05151   -1.409 0.158891
## GenrePlatform   0.34567    0.06600    5.237 1.65e-07 ***
## GenrePuzzle    -0.08888    0.07829   -1.135 0.256292
## GenreRacing     0.01634    0.05785    0.282 0.777645
## GenreRole-Playing 0.10233    0.05435    1.883 0.059733 .
## GenreShooter    0.25353    0.05669    4.472 7.80e-06 ***
## GenreSimulation -0.09978    0.06589   -1.514 0.129972
## GenreSports     0.06256    0.04715    1.327 0.184597
## GenreStrategy  -0.26765    0.07313   -3.660 0.000254 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.557 on 13264 degrees of freedom
## Multiple R-squared:  0.01074,    Adjusted R-squared:  0.009917
## F-statistic: 13.09 on 11 and 13264 DF,  p-value: < 2.2e-16
```

When the `summary()` function was called, it displayed what the call was for the linear model, residuals, coefficients, as well as a few other value pertaining to the target and predictor. Information within the residuals display a bit of data exploration statistics including the min, max, median, and the first and third quarters.

In coefficients are the intercept as well as predictors from the Genre column. An estimated coefficient, standard error, t-value, and p-value are also shown. There are asterisks next to Adventure, Platform, Shooter, and Strategy which suggests that those four genres may be good predictors of Global\_Sales. Standard error is a variation of the estimate given for each predictor and the intercept. The t-value signifies how little of a relationship two variables would have, or how true the null hypothesis is. However, the p-value tells the opposite. With a low p-value, the null hypothesis can be rejected meaning that there is a potential relationship between the two variables.

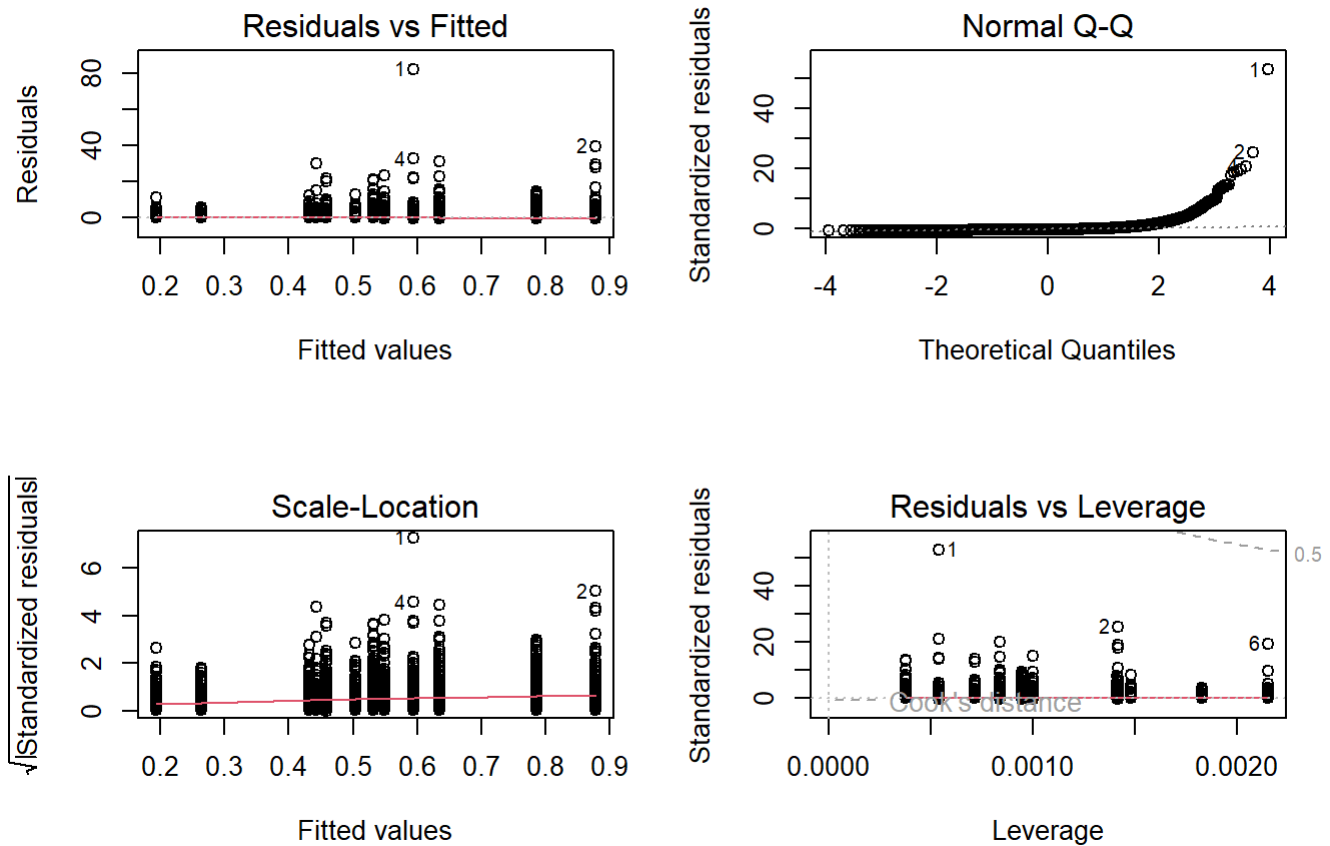
The last part of summary are values that portray how well each coefficient modeled the true data. It seems that the model is off by about 1.6 million copies. An R squared value is also given, which is 0.009917. This implies that Genre most likely does not have a relationship with Global\_Sales. However, the F-statistic (accounts for all predictors) is higher than 11 and the p-value is low, so this model may display some kind of significance.

## Plotting the residuals

The following are plotting the residuals generated from `lm()`.

```
par(mfrow=c(2,2))
plot(lm1)
```





Starting with the Residuals vs Fitted plot, there technically is a linear pattern. However, the residuals are not spread out evenly and well, but rather clump up in columns at particular values. This may mean that a non-linear pattern can be present. The next plot is the Normal Q-Q which signifies if the residuals follow a straight line well. It seems that way at first, until around the second quantile. The residuals start to significantly deviate from the line which is not ideal. So, the residuals may not be normally distributed. The Scale-Location plot is used to check if residuals are spread equally within the range of predictors. However, like the first plot, the problem lies with how the residuals stack in columns. This could indicate that the predictor of genre may not predict this well. Lastly is the residuals vs leverage plot which indicates if any extremities could influence the regression line. Oddly enough, in this case, none of the residuals appears outside Cook's distance lines. This implies that none of the residuals would affect how the linear regression would be formed.

## Multiple linear regression

Multiple linear regression is carried out in an attempt to make the data better. Two predictors are used instead by including Platform into the model. In doing so, warnings are generated about leverages being one. However, these warnings are ignored.

```
lm2 <- lm(Global_Sales~Genre+Platform, data=train)
summary(lm2)
```

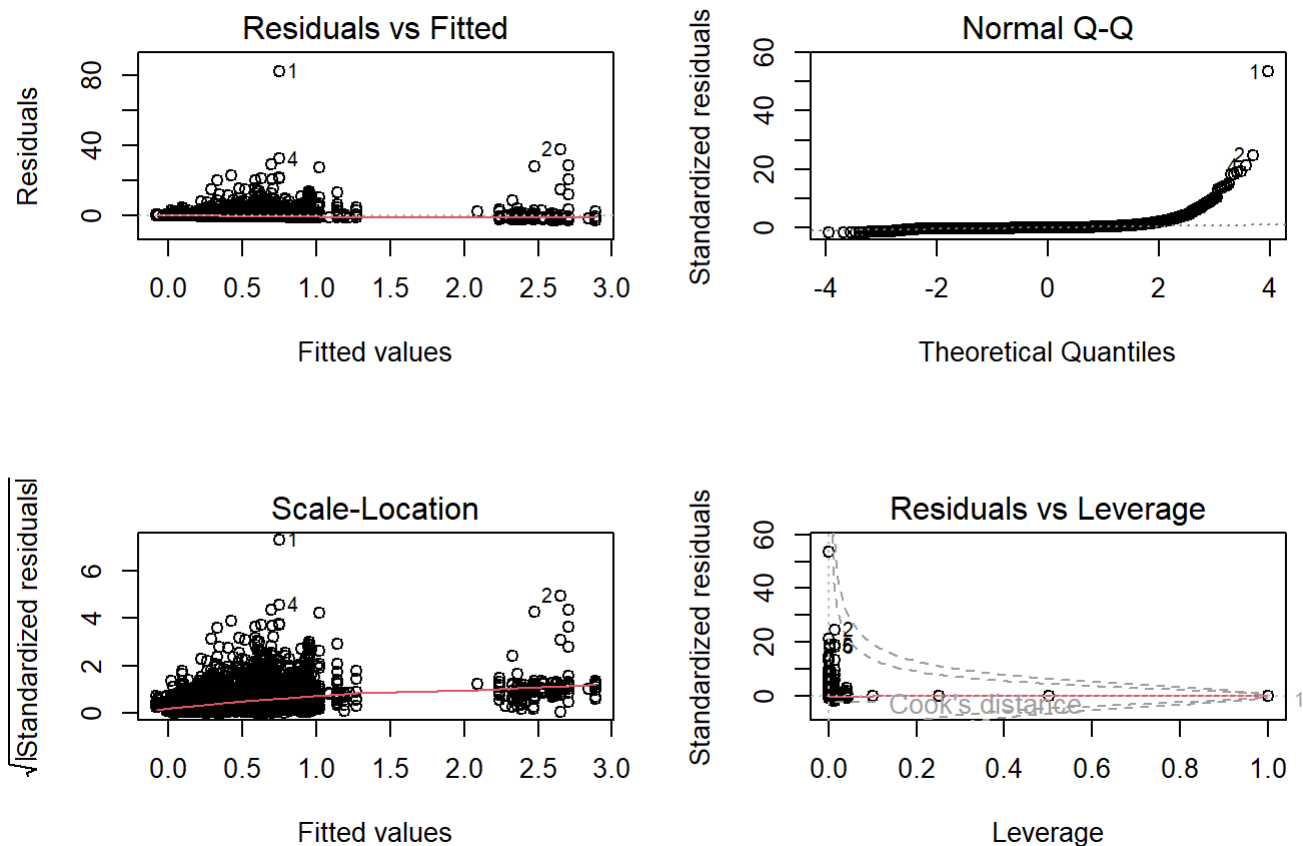
```
##
## Call:
## lm(formula = Global_Sales ~ Genre + Platform, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.818 -0.453 -0.240  0.014  81.988
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.69666    0.15252   4.568 4.97e-06 ***
## GenreAdventure -0.26986    0.05704  -4.731 2.26e-06 ***
## GenreFighting  -0.02836    0.06735  -0.421  0.67369
## GenreMisc      -0.07564    0.05158  -1.467  0.14250
## GenrePlatform   0.29143    0.06641   4.389 1.15e-05 ***
## GenrePuzzle    -0.11820    0.07890  -1.498  0.13412
## GenreRacing     0.02011    0.05780   0.348  0.72797
## GenreRole-Playing 0.11139    0.05421   2.055  0.03990 *
## GenreShooter    0.24288    0.05659   4.292 1.78e-05 ***
## GenreSimulation -0.06234    0.06615  -0.942  0.34598
## GenreSports     0.02670    0.04720   0.566  0.57163
## GenreStrategy  -0.19144    0.07379  -2.594  0.00949 **
## Platform3D0     -0.46263    1.09500  -0.422  0.67267
## Platform3DS     -0.19659    0.16914  -1.162  0.24514
## PlatformDC      -0.36316    0.28398  -1.279  0.20098
## PlatformDS      -0.29076    0.15630  -1.860  0.06288 .
## PlatformGB      1.90021    0.22935   8.285 < 2e-16 ***
## PlatformGBA     -0.32820    0.16317  -2.011  0.04430 *
## PlatformGC      -0.39778    0.16826  -2.364  0.01809 *
## PlatformGEN      0.28180    0.34814   0.809  0.41827
## PlatformGG      -0.94809    1.54114  -0.615  0.53844
## PlatformN64     -0.02068    0.17948  -0.115  0.90828
## PlatformNES      1.66099    0.22907   7.251 4.37e-13 ***
## PlatformNG      -0.56380    0.51070  -1.104  0.26962
## PlatformPC      -0.41029    0.16183  -2.535  0.01125 *
## PlatformPCFX    -0.77805    1.54080  -0.505  0.61359
## PlatformPS      -0.09624    0.15960  -0.603  0.54651
## PlatformPS2     -0.11348    0.15621  -0.726  0.46755
## PlatformPS3      0.00384    0.15863   0.024  0.98069
## PlatformPS4      0.19940    0.17875   1.116  0.26464
## PlatformPSP     -0.43679    0.15978  -2.734  0.00627 **
## PlatformPSV     -0.51230    0.17331  -2.956  0.00312 **
## PlatformSAT     -0.49615    0.19964  -2.485  0.01296 *
## PlatformSCD     -0.57102    1.54066  -0.371  0.71092
## PlatformSNES    -0.02773    0.18837  -0.147  0.88294
## PlatformTG16    -0.28680    1.54090  -0.186  0.85235
## PlatformWii      0.02881    0.15884   0.181  0.85606
## PlatformWiiU    -0.12558    0.20817  -0.603  0.54633
## PlatformWS      -0.47414    0.78232  -0.606  0.54448
## PlatformX360     0.00931    0.15871   0.059  0.95322
## PlatformXB      -0.44202    0.16314  -2.709  0.00675 **
## PlatformXOne    -0.05645    0.19162  -0.295  0.76831
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.533 on 13234 degrees of freedom
## Multiple R-squared:  0.04321,    Adjusted R-squared:  0.04025
## F-statistic: 14.58 on 41 and 13234 DF,  p-value: < 2.2e-16
```

```
par(mfrow=c(2,2))
plot(lm2)
```

```
## Warning: not plotting observations with leverage one:
##    2214, 2767
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```



## Linear regression improvement attempt

In a third attempt to better the data, JP\_Sales was used instead as there seems to be less outliers associated with it. This was in an attempt to check if the Global\_Sales was a problem. Warnings were also issued and ignored.

```
lm3 <- lm(JP_Sales~Platform+Genre, data=train) #Tried an interaction between platform and Genre  
summary(lm3)
```





```
##
## Call:
## lm(formula = JP_Sales ~ Platform + Genre, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.0969 -0.0631 -0.0376  0.0078  9.2587
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.0010110   0.0282302    0.036  0.97143
## Platform3D0     0.0491214   0.2026798    0.242  0.80851
## Platform3DS     0.1707898   0.0313064    5.455 4.97e-08 ***
## PlatformDC      0.1484749   0.0525630    2.825  0.00474 **
## PlatformDS      0.0623856   0.0289311    2.156  0.03107 *
## PlatformGB      0.8085782   0.0424515   19.047 < 2e-16 ***
## PlatformGBA     0.0439586   0.0302015    1.456  0.14555
## PlatformGC      0.0240793   0.0311437    0.773  0.43944
## PlatformGEN     0.0740447   0.0644388    1.149  0.25055
## PlatformGG      0.0120412   0.2852583    0.042  0.96633
## PlatformN64     0.1036536   0.0332204    3.120  0.00181 **
## PlatformNES     1.0542246   0.0424005   24.864 < 2e-16 ***
## PlatformNG      0.0849948   0.0945279    0.899  0.36859
## PlatformPC     -0.0196797   0.0299540   -0.657  0.51119
## PlatformPCFX   -0.1226870   0.2851960   -0.430  0.66707
## PlatformPS      0.0965209   0.0295421    3.267  0.00109 **
## PlatformPS2     0.0498715   0.0289133    1.725  0.08458 .
## PlatformPS3     0.0469945   0.0293626    1.600  0.10951
## PlatformPS4     0.0249116   0.0330851    0.753  0.45149
## PlatformPSP     0.0313156   0.0295744    1.059  0.28968
## PlatformPSV     0.0211725   0.0320784    0.660  0.50925
## PlatformSAT     0.1594190   0.0369532    4.314 1.61e-05 ***
## PlatformSCD     0.0392584   0.2851701    0.138  0.89051
## PlatformSNES    0.4130571   0.0348658   11.847 < 2e-16 ***
## PlatformTG16    0.1503946   0.2852149    0.527  0.59799
## PlatformWii     0.0479449   0.0294004    1.631  0.10296
## PlatformWiiU    0.0899884   0.0385308    2.335  0.01953 *
## PlatformWS      0.0976776   0.1448036    0.675  0.49997
## PlatformX360   -0.0006301   0.0293757   -0.021  0.98289
## PlatformXB     -0.0036254   0.0301966   -0.120  0.90444
## PlatformXOne   -0.0075311   0.0354688   -0.212  0.83185
## GenreAdventure -0.0114056   0.0105582   -1.080  0.28005
## GenreFighting   0.0270406   0.0124670    2.169  0.03010 *
## GenreMisc      -0.0002694   0.0095467   -0.028  0.97749
## GenrePlatform   0.0269478   0.0122912    2.192  0.02837 *
## GenrePuzzle     -0.0088593   0.0146031   -0.607  0.54408
## GenreRacing     -0.0082239   0.0106987   -0.769  0.44210
## GenreRole-Playing 0.1516760   0.0100334   15.117 < 2e-16 ***
## GenreShooter    -0.0089410   0.0104745   -0.854  0.39334
## GenreSimulation  0.0184207   0.0122434    1.505  0.13247
## GenreSports     -0.0034230   0.0087371   -0.392  0.69523
## GenreStrategy   0.0159468   0.0136583    1.168  0.24301
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2837 on 13234 degrees of freedom
## Multiple R-squared:  0.163, Adjusted R-squared:  0.1604
## F-statistic: 62.88 on 41 and 13234 DF, p-value: < 2.2e-16
```

```
par(mfrow=c(2,2))
plot(lm3)
```

```
## Warning: not plotting observations with leverage one:
## 2767, 6391
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```



## Results

Between the three models, models two and three looks better in terms of the first and third plot. Unlike the first plot, the residuals are less consolidated and more spread out. At the same time though, it is not a drastic improvement over plot one as there is still some clumping in the lower and upper end of the fitted values for plot two. In plot three it is just the lower end. I believe before any more analysis would be done, the third model could

be eliminated just because the data did not seem to be any better for JP\_Sales over Global\_Sales in the second model. I believe the second model would be better over the first as there is more data to work with in the first and third plots. However, there is a bit of concern with plot two as the lower end does not fit the linear line. Even more concern is expressed with the fourth plot since the residual 1 is at the tip of the Cook's distance, meaning that there could be something affecting the linear regression.

## Correlation and MSE

### Model 1 Evaluation

```
pred1 <- predict(lm1, newdata=test)
cor1 <- cor(pred1, test$Global_Sales)
mse1 <- mean((pred1-test$Global_Sales)^2)
rmse1 <- sqrt(mse1)

print(paste('correlation:', cor1))
```

```
## [1] "correlation: 0.127974758171416"
```

```
print(paste('mse:', mse1))
```

```
## [1] "mse: 2.265806009761"
```

```
print(paste('rmse:', rmse1))
```

```
## [1] "rmse: 1.50525944931796"
```

### Model 2 Evaluation

```
pred2 <- predict(lm2, newdata=test)
cor2 <- cor(pred2, test$Global_Sales)
mse2 <- mean((pred2-test$Global_Sales)^2)
rmse2 <- sqrt(mse2)

print(paste('correlation:', cor2))
```

```
## [1] "correlation: 0.218379234764752"
```

```
print(paste('mse:', mse2))
```

```
## [1] "mse: 2.19279293135131"
```

```
print(paste('rmse:', rmse2))
```

```
## [1] "rmse: 1.48080820208132"
```

### Model 3 Evaluation

```
pred3 <- predict(lm3, newdata=test)
cor3 <- cor(pred3, test$Global_Sales)
mse3 <- mean((pred3-test$Global_Sales)^2)
rmse3 <- sqrt(mse3)

print(paste('correlation:', cor3))
```

```
## [1] "correlation: 0.152110255322097"
```

```
print(paste('mse:', mse3))
```

```
## [1] "mse: 2.48497871004454"
```

```
print(paste('rmse:', rmse3))
```

```
## [1] "rmse: 1.57638152426516"
```

Between all three models, it seems like model two has the highest correlation, but it is still a bad correlation being only approximately 0.22. Despite being a poor model, model two is still the best as comparing the mse and rmse values with model one, they are both slightly lower. This implies that model two is just a slightly better fit of a model for Global\_Sales. I think this may have occurred because with including platform as another predictor in model two, it gave a few more predictors that were good for the model thus helping to support correlation to Global\_Sales.